

Longitudinal Analysis

Richard White

2018-04-12

Contents

1	Syllabus	5
2	Reference	7
3	Variables	9
3.1	Showing that <code>tscount::tsglm</code> gets the same results as <code>MASS::glmPQL</code>	18
4	Variables	43
4.1	Showing that <code>tscount::tsglm</code> gets the same results as <code>MASS::glmPQL</code>	50
5	Variables	75
6	Variables	77
7	Variables	83

Chapter 1

Syllabus

Instructor: Richard White [richard.white@fhi.no]

Time: 09:30 - 15:00, 18th September 2017

Location: Main auditorium, L8, Lindern Campus, Folkehelseinstituttet, Oslo

Language: English

Format and Procedures

09:00 - 10:00: Lecture 1

10:00 - 10:10: Break

10:10 - 11:10: Lecture 2

10:10 - 10:15: Break

11:15 - 11:45: Examples from FHI

Description

This course will provide a basic overview of general statistical methodology that can be useful in the areas of infectious diseases, environmental medicine, and labwork. By the end of this course, students will be able to identify appropriate statistical methods for a variety of circumstances.

This course will **not** teach students how to implement these statistical methods, as there is not sufficient time. The aim of this course is to enable the student to identify which methods are required for their study, allowing the student to identify their needs for subsequent methods courses, self-learning, or external help.

You should register for this course if you are one of the following:

- Have experience with applying statistical methods, but are sometimes confused or uncertain as to whether or not you have selected the correct method.
- Do not have experience with applying statistical methods, and would like to get an overview over which methods are applicable for your projects so that you can then undertake further studies in these areas.

Lecture 1

1. Identifying continuous, categorical, count, and censored variables
2. Identifying exposure and outcome variables
3. Identifying when t-tests (paired and unpaired) should be used
4. Identifying when non-parametric t-test equivalents should be used
5. Identifying when ANOVA should be used
6. Identifying when linear regression should be used

7. Identifying the similarities between t-tests, ANOVA, and regression
8. Identifying when logistic regression models should be used
9. Identifying when Poisson/negative binomial and cox regression models should be used
10. Identifying when chi-squared/fisher's exact test should be used

Lecture 2

1. Identifying when data does not have any dependencies (i.e. all observations are independent of each other) versus when data has complicated dependencies (i.e. longitudinal data, matched data, multiple cohorts)
2. Identifying when mixed effects regression models should be used
3. Identifying when conditional logistic regression models should be used
4. (TBD) Understanding the different imputation methods used when lab data is below the limit of detection (LOD)
5. (TBD) Understanding the best practices for data files and project folders

Prerequisites

To participate in this course it is recommended that you have some experience with either research or data.

Additional information

For the last 30 minutes of the course we will be going through examples of analyses performed at FHI and identifying which statistical methods are appropriate. If you would like your analysis to be featured/included in this section, please send an email to richard.white@fhi.no briefly describing your problem.

Chapter 2

Reference

AIM														
Hypothesis Testing	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes
Effect estimation	No	No	No	No	No	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes
OUTCOME														
Continuous	No	Yes	Yes	Yes	Yes	Yes	No	No	No	No	Yes	No	No	No
Binary	No	No	No	No	No	No	Yes	No	No	No	No	Yes	No	No
Categorical	Yes	No	No	No	No	No	No	No	No	No	No	No	No	No
Censored	No	No	No	No	No	No	No	No	Yes	No	No	No	No	No
Count	No	No	No	No	No	No	No	Yes	Yes	No	No	No	Yes	Yes
EXPOSURE														
Continuous	No	No	No	No	No	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes
Binary	No	No	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes
Categorical	Yes	No	No	No	No	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes
Censored	No	No	No	No	No	No	No	No	No	No	No	No	No	No
Count	No	No	No	No	No	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes
DEPENDENCIES														
Data	No	No	Yes	No	No	No	No	No	No	No	Yes	Yes	Yes	Yes
	Chi-Squared	One Sample T-Test	Two Sample Paired T-Test	Two Sample Unpaired T-Test	ANOVA	Linear regression	Logistic regression	Poisson regression	Negative-binomial regression	Cox regression	Mixed effects linear regression	Mixed effects logistic regression	Mixed effects Poisson regression	Mixed effects negative-binomial regression

Yes
No

Chapter 3

Variables

```
library(data.table)
library(ggplot2)
set.seed(4)


AMPLITUDE <- 1.5
SEASONAL_HORIZONTAL_SHIFT <- 20

d <- data.table(date=seq.Date(
  from=as.Date("2000-01-01"),
  to=as.Date("2018-12-31"),
  by=1))
d[,year:=as.numeric(format.Date(date,"%G"))]
d[,week:=as.numeric(format.Date(date,"%V"))]
d[,month:=as.numeric(format.Date(date,"%m"))]
d[,yearMinus2000:=year-2000]

d[,dayOfYear:=as.numeric(format.Date(date,"%j"))]
d[,seasonalEffect:=sin(2*pi*(dayOfYear-SEASONAL_HORIZONTAL_SHIFT)/365)]
d[,mu := exp(0.1 + yearMinus2000*0.1 + seasonalEffect*AMPLITUDE)]
d[,y:=rpois(.N,mu)]
```

Showing the true data

```
q <- ggplot(d,aes(x=date))
q <- q + geom_point(mapping=aes(y=y))
q <- q + geom_line(mapping=aes(y=mu),colour="red")
q
```




Investigating the data

We take a quick look, but don't see much

```
q <- ggplot(d, aes(x=date, y=y))
q <- q + geom_point()
q <- q + stat_smooth(colour="red")
q
```


```
## `geom_smooth()` using method = 'gam'
```



We then drill down into a few years, and see a clear seasonal trend

```
q <- ggplot(d[year %in% c(2005:2010)], aes(x=dayOfYear, y=y))
q <- q + facet_wrap(~year)
q <- q + geom_point()
q <- q + stat_smooth(colour="red")
q
```


```
## `geom_smooth()` using method = 'loess'
```



The Lomb-Scargle Periodogram shows a clear seasonality with a period of 365 days

```
lomb::lsp(d$y, from=100, to=500, ofac=1, type="period")
```

Lomb–Scargle Periodogram



We then generate two new variables `cos365` and `sin365` and perform a simple poisson regression:

```
d[,cos365:=cos(dayOfYear*2*pi/365)]
d[,sin365:=sin(dayOfYear*2*pi/365)]

fit0 <- glm(y~yearMinus2000, data=d, family=poisson())
fit1 <- glm(y~yearMinus2000+sin365 + cos365, data=d, family=poisson())

print(lmtest::lrtest(fit0, fit1))

## Likelihood ratio test
##
## Model 1: y ~ yearMinus2000
## Model 2: y ~ yearMinus2000 + sin365 + cos365
##   #Df LogLik Df Chisq Pr(>Chisq)
## 1    2 -27287
## 2    4 -12805  2 28963 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

print(summary(fit1))

##
## Call:
## glm(formula = y ~ yearMinus2000 + sin365 + cos365, family = poisson(),
##      data = d)
##
##
## Deviance Residuals:
```

```
##      Min      1Q   Median      3Q      Max
## -3.7499 -0.9167 -0.1370  0.5955  3.2193
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.086654  0.014940   5.80 6.62e-09 ***
## yearMinus2000 0.100461  0.001049  95.75 < 2e-16 ***
## sin365       1.428417  0.010434 136.90 < 2e-16 ***
## cos365      -0.512912  0.008666 -59.19 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for poisson family taken to be 1)
##
##      Null deviance: 46221.4  on 6939  degrees of freedom
## Residual deviance:  7259.2  on 6936  degrees of freedom
## AIC: 25619
##
## Number of Fisher Scoring iterations: 5
```


We see a clear significant seasonal effect. We can then use trigonometry to back-calculate the `cos365` and `sin365` variables to amplitude and location of peak/troughs:

```
b1 <- 1.428417 # sin coefficient
b2 <- -0.512912 # cos coefficient
amplitude <- sqrt(b1^2 + b2^2)
p <- atan(b1/b2) * 365/2/pi
if (p > 0) {
  peak <- p
  trough <- p + 365/2
} else {
  peak <- p + 365/2
  trough <- p + 365
}
if (b1 < 0) {
  g <- peak
  peak <- trough
  trough <- g
}
print(sprintf("amplitude is estimated as %s, peak is estimated as %s, trough is estimated as %s",round(
## [1] "amplitude is estimated as 1.52, peak is estimated as 111, trough is estimated as 294"
print(sprintf("true values are: amplitude: %s, peak: %s, trough: %s",round(AMPLITUDE,2),round(365/4+SEA
## [1] "true values are: amplitude: 1.5, peak: 111, trough: 294"
```


We now investigate our residuals to determine if we have a good fit:

```
d[,residuals:=residuals(fit1, type = "response")]
d[,predicted:=predict(fit1, type = "response")]
q <- ggplot(d,aes(x=predicted,y=residuals))
q <- q + geom_point()
q <- q + stat_smooth(colour="red")
q

## `geom_smooth()` using method = 'gam'
```




```
# this is for AR  
pacf(d$residuals)
```



```
# this is for MA  
acf(d$residuals)
```


Series d\$residuals



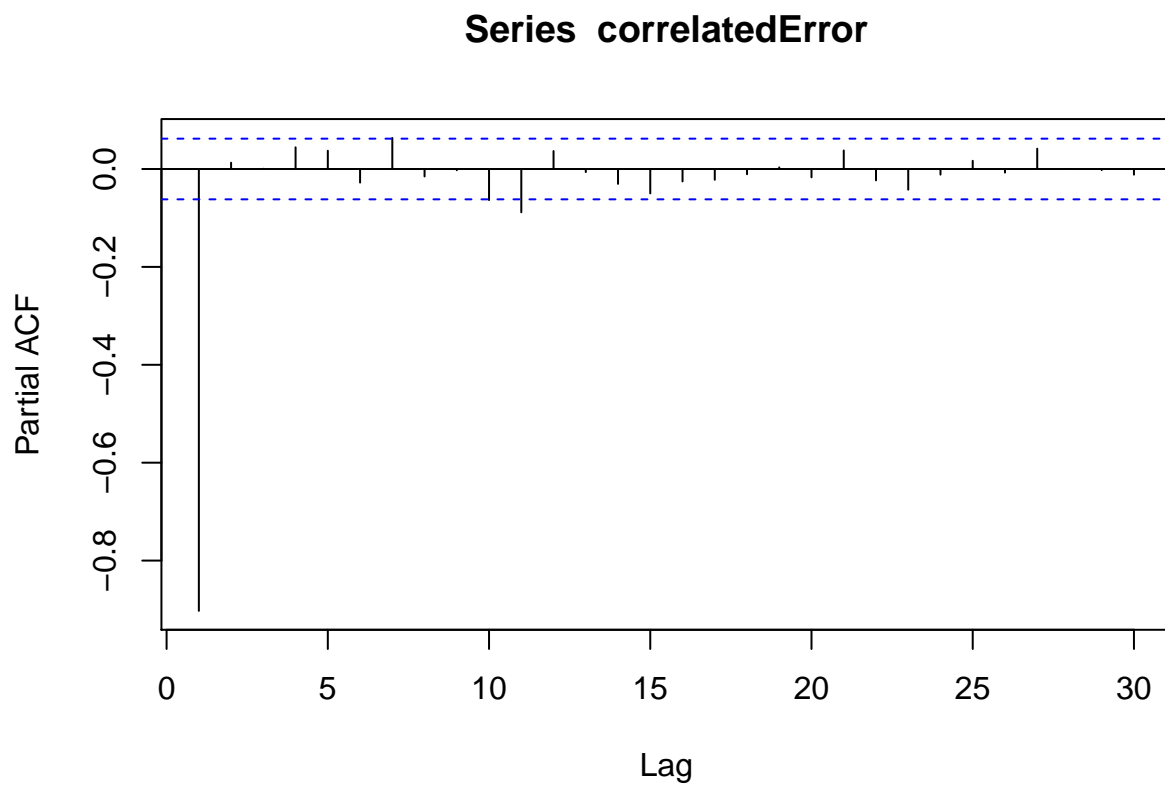
We see a clear significant seasonal effect. We can then use trigonometry to back-calculate the `cos365` and `sin365` variables to amplitude and location of peak/troughs:

```
b1 <- 0.1934 # sin coefficient
b2 <- 0.1018 # cos coefficient
amplitude <- sqrt(b1^2 + b2^2)
p <- atan(b1/b2) * 365/2/pi
if (p > 0) {
  peak <- p
  trough <- p + 365/2
} else {
  peak <- p + 365/2
  trough <- p + 365
}
if (b1 < 0) {
  g <- peak
  peak <- trough
  trough <- g
}
print(sprintf("amplitude is %s, peak is at %s, trough is at %s",round(amplitude,2),round(peak),round(trough)))
```


```
## [1] "amplitude is 0.22, peak is at 63, trough is at 246"
```

3.1 Showing that `tscount::tsglm` gets the same results as `MASS::glmPQL`

```
library(MASS)
correlatedError <- as.numeric(arima.sim(model=list("ar"=c(-0.9)), n=1000, rand.gen = rnorm))
pacf(correlatedError) # this is for AR
```



```
acf(correlatedError) # this is for MA
```




```
d <- data.frame(correlatedError)
d$independentError <- rnorm(nrow(d))
d$x <- rnorm(nrow(d))
d$yCorrelated <- 2*d$x+d$correlatedError
d$yIndependent <- 2*d$x+d$independentError
d$ID <- 1
d$time <- 1:nrow(d)
```

```
summary(lm(yIndependent~x,data=d))
```


```
##
## Call:
## lm(formula = yIndependent ~ x, data = d)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3.2196 -0.7132  0.0044  0.6760  2.9593
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.06795   0.03144    2.161  0.0309 *
## x            1.97958   0.03159   62.657 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.9943 on 998 degrees of freedom
```

```
## Multiple R-squared:  0.7973, Adjusted R-squared:  0.7971
## F-statistic:  3926 on 1 and 998 DF,  p-value: < 2.2e-16
summary(fit <- lm(yCorrelated~x,data=d))

##
## Call:
## lm(formula = yCorrelated ~ x, data = d)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -7.1132 -1.6599  0.0492  1.6199  6.5856
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.0008572  0.0741155  -0.012   0.991
## x            2.0247397  0.0744706  27.188 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.344 on 998 degrees of freedom
## Multiple R-squared:  0.4255, Adjusted R-squared:  0.4249
## F-statistic: 739.2 on 1 and 998 DF,  p-value: < 2.2e-16
pacf(residuals(fit)) # this is for AR
```



```
acf(residuals(fit)) # this is for MA
```




```
# independent data, no correlation structure needed
fit <- MASS::glmmpQL(yIndependent ~ x, random = ~ 1 | ID,
  family = gaussian, data = d,
  correlation=nlme::corAR1())
```

```
## iteration 1
```


```
summary(fit)
```

```
## Linear mixed-effects model fit by maximum likelihood
## Data: d
##   AIC BIC logLik
##   NA  NA   NA
##
## Random effects:
## Formula: ~1 | ID
##      (Intercept) Residual
## StdDev: 3.157677e-05 0.9933179
##
## Correlation Structure: AR(1)
## Formula: ~1 | ID
## Parameter estimate(s):
##      Phi
## 0.05315098
## Variance function:
## Structure: fixed weights
## Formula: ~invwt
```

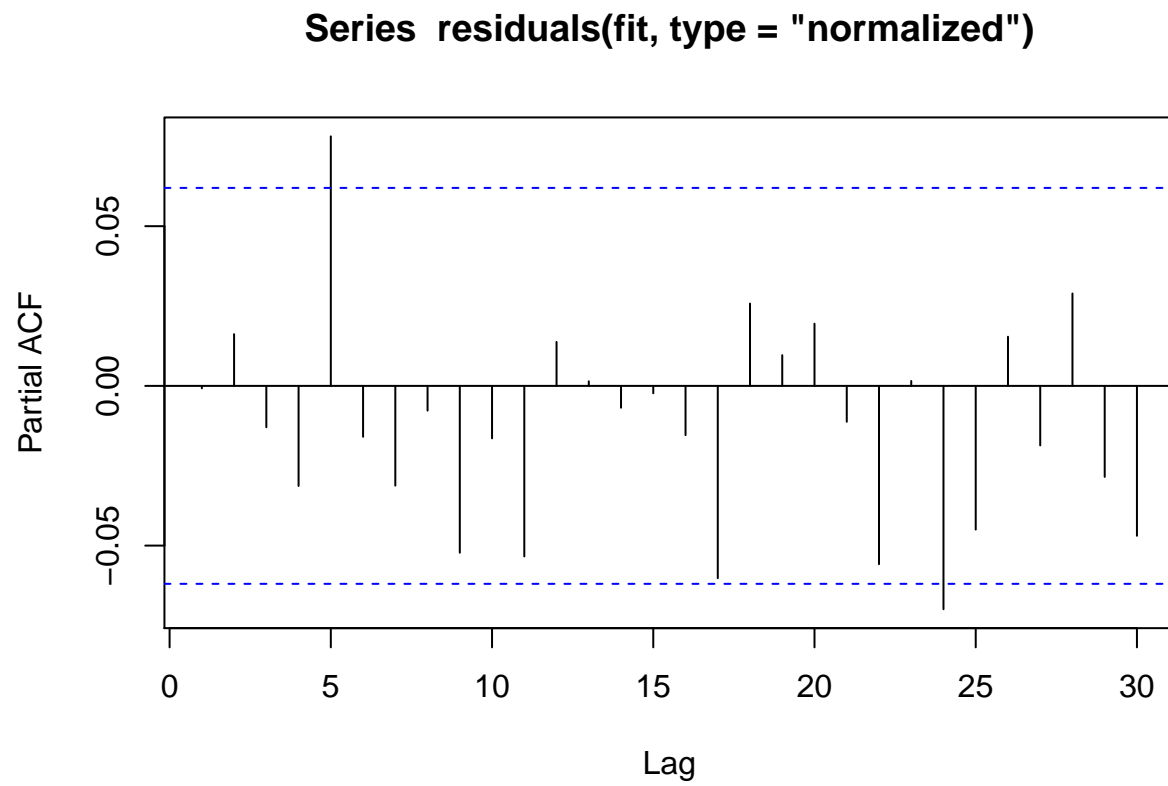
```
## Fixed effects: yIndependent ~ x
##               Value Std.Error DF  t-value p-value
## (Intercept) 0.0679433 0.03315919 998  2.04900  0.0407
## x           1.9792647 0.03162358 998 62.58825  0.0000
## Correlation:
##   (Intr)
## x 0.001
##
## Standardized Within-Group Residuals:
##      Min      Q1      Med      Q3      Max
## -3.241717414 -0.718010178  0.003911461  0.680810722  2.978619340
##
## Number of Observations: 1000
## Number of Groups: 1
pacf(residuals(fit, type = "response")) # this is for AR
```




```
acf(residuals(fit, type = "response")) # this is for MA
```



```
pacf(residuals(fit, type = "normalized")) # this is for AR
```



```
acf(residuals(fit, type = "normalized")) # this is for MA
```


Series residuals(fit, type = "normalized")


```
# dependent data, needs correlation structure, no correlation structure
fit <- MASS::glmmpQL(yCorrelated ~ x, random = ~ 1 | ID,
  family = gaussian, data = d)
```

```
## iteration 1
```


```
summary(fit)
```

```
## Linear mixed-effects model fit by maximum likelihood
## Data: d
##   AIC BIC logLik
##   NA  NA    NA
##
## Random effects:
## Formula: ~1 | ID
##      (Intercept) Residual
## StdDev: 7.66613e-05 2.341391
##
## Variance function:
## Structure: fixed weights
## Formula: ~invwt
## Fixed effects: yCorrelated ~ x
##              Value Std.Error DF   t-value p-value
## (Intercept) -0.0008572 0.07411551 998 -0.011566  0.9908
## x              2.0247397 0.07447061 998 27.188441  0.0000
## Correlation:
##   (Intr)
```

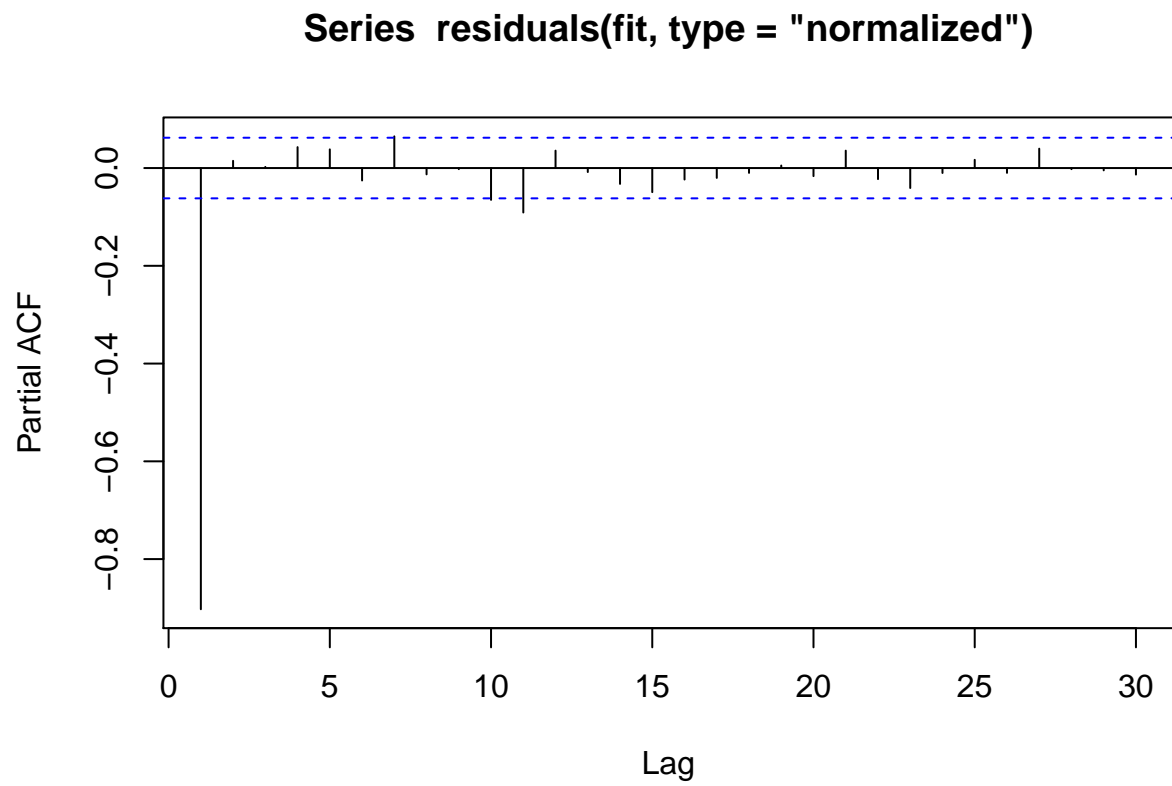
```
## x 0.001
##
## Standardized Within-Group Residuals:
##      Min      Q1      Med      Q3      Max
## -3.03803335 -0.70892801  0.02100393  0.69187212  2.81270067
##
## Number of Observations: 1000
## Number of Groups: 1
pacf(residuals(fit, type = "response")) # this is for AR
```




```
acf(residuals(fit, type = "response")) # this is for MA
```



```
pacf(residuals(fit, type = "normalized")) # this is for AR
```



```
acf(residuals(fit, type = "normalized")) # this is for MA
```



```
# dependent data, correct correlation structure
fit <- MASS::glmmpQL(yCorrelated ~ x, random = ~ 1 | ID,
  family = gaussian, data = d,
  correlation=nlme::corAR1())
```


```
## iteration 1
```

```
## iteration 2
```


```
summary(fit)
```

```
## Linear mixed-effects model fit by maximum likelihood
## Data: d
## AIC BIC logLik
## NA NA NA
##
## Random effects:
## Formula: ~1 | ID
## (Intercept) Residual
## StdDev: 1.698085e-05 2.341112
##
## Correlation Structure: AR(1)
## Formula: ~1 | ID
## Parameter estimate(s):
## Phi
## -0.9035598
## Variance function:
## Structure: fixed weights
```


```
## Formula: ~invwt
## Fixed effects: yCorrelated ~ x
##           Value Std.Error DF t-value p-value
## (Intercept) 0.0010465 0.01668822 998 0.06271 0.95
## x           2.0007666 0.02289286 998 87.39695 0.00
## Correlation:
## (Intr)
## x 0.002
##
## Standardized Within-Group Residuals:
##      Min      Q1      Med      Q3      Max
## -3.05408286 -0.70634334 0.02403998 0.69306761 2.81549938
##
## Number of Observations: 1000
## Number of Groups: 1
pacf(residuals(fit, type = "response")) # this is for AR
```




```
acf(residuals(fit, type = "response")) # this is for MA
```



```
pacf(residuals(fit, type = "normalized")) # this is for AR
```




```
acf(residuals(fit, type = "normalized")) # this is for MA
```


```
library(MASS)

bacteria$x <- 1
fit <- glmmpQL(as.numeric(y) ~ trt + I(week > 2), random = ~ 1 | x,
               family = poisson, data = bacteria)
```

```
## iteration 1
acf(residuals(fit,type="normalized"))
```




```
pacf(residuals(fit,type="normalized"))
```




```
bacteria$x <- 1
fit <- glmmPQL(as.numeric(y) ~ trt + I(week > 2), random = ~ 1 | x,
               family = poisson, data = bacteria,
               correlation=nlme::corAR1())
```

```
## iteration 1
```

```
## iteration 2
acf(residuals(fit,type="normalized"))
```




```
pacf(residuals(fit,type="normalized"))
```




```
fit <- glmmPQL(as.numeric(y) ~ trt + I(week > 2), random = ~ 1 | ID,
              family = poisson, data = bacteria)
```

```
## iteration 1
## iteration 2
```

```
## iteration 3
acf(residuals(fit,type="normalized"))
```



```
pacf(residuals(fit,type="normalized"))
```



```
fit <- glmmPQL(as.numeric(y) ~ trt + I(week > 2), random = ~ 1 | ID,
               family = poisson, data = bacteria,
               correlation=nlme::corAR1(form=~ 1 | ID))
```


```
## iteration 1
```

```
## iteration 2
```


```
## iteration 3
```

```
## iteration 4
```

```
acf(residuals(fit,type="normalized"))
```



```
pacf(residuals(fit,type="normalized"))
```

Chapter 4

Variables

```
library(data.table)
library(ggplot2)
set.seed(4)


AMPLITUDE <- 1.5
SEASONAL_HORIZONTAL_SHIFT <- 20

d <- data.table(date=seq.Date(
  from=as.Date("2000-01-01"),
  to=as.Date("2018-12-31"),
  by=1))
d[,year:=as.numeric(format.Date(date,"%G"))]
d[,week:=as.numeric(format.Date(date,"%V"))]
d[,month:=as.numeric(format.Date(date,"%m"))]
d[,yearMinus2000:=year-2000]
d[,dayOfSeries:=1:.N]

d[,dayOfYear:=as.numeric(format.Date(date,"%j"))]
d[,seasonalEffect:=sin(2*pi*(dayOfYear-SEASONAL_HORIZONTAL_SHIFT)/365)]
d[,mu := exp(0.1 + yearMinus2000*0.1 + seasonalEffect*AMPLITUDE)]
d[,y:=rpois(.N,mu)]
d[,y:=round(as.numeric(arima.sim(model=list("ar"=c(0.5)), rand.gen = rpois, n=nrow(d), lambda=mu)))]

q <- ggplot(d[year %in% c(2005:2010)],aes(x=dayOfYear,y=y))
q <- q + facet_wrap(~year)
q <- q + geom_point()
q <- q + stat_smooth(colour="red")
q
```


```
## `geom_smooth()` using method = 'loess'
```



The Lomb-Scargle Periodogram shows a clear seasonality with a period of 365 days

```
lomb::lsp(d$y, from=50, to=500, ofac=1, type="period")
```

Lomb–Scargle Periodogram



```
d[,cos365:=cos(dayOfYear*2*pi/365)]
d[,sin365:=sin(dayOfYear*2*pi/365)]

fit0 <- glm(y~yearMinus2000, data=d, family=poisson())
fit1 <- glm(y~yearMinus2000+sin365 + cos365, data=d, family=poisson())

print(lmtest::lrtest(fit0, fit1))
```

```
## Likelihood ratio test
##
## Model 1: y ~ yearMinus2000
## Model 2: y ~ yearMinus2000 + sin365 + cos365
##   #Df LogLik Df Chisq Pr(>Chisq)
## 1    2 -43124
## 2    4 -14542  2 57163 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```


```
print(summary(fit1))
```

```
##
## Call:
## glm(formula = y ~ yearMinus2000 + sin365 + cos365, family = poisson(),
##      data = d)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
```

```
## -2.6774 -0.6738 -0.0503 0.4920 3.5820
##
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.7981246  0.0105300   75.80  <2e-16 ***
## yearMinus2000 0.0991480  0.0007416  133.70  <2e-16 ***
## sin365       1.4074818  0.0073418  191.71  <2e-16 ***
## cos365      -0.5390314  0.0061513  -87.63  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for poisson family taken to be 1)
##
## Null deviance: 81832.6 on 6939 degrees of freedom
## Residual deviance: 5217.8 on 6936 degrees of freedom
## AIC: 29093
##
## Number of Fisher Scoring iterations: 4
d[,residuals:=residuals(fit1, type = "response")]
d[,predicted:=predict(fit1, type = "response")]


# this is for AR
pacf(d$residuals)
```

Series d\$residuals



```
# this is for MA
acf(d$residuals)
```

Series d\$residuals



This means our model is bad, we have autocorrelation.

```
d[,ID:=1]
# this is for MA
fit <- MASS::glmPQL(y~yearMinus2000+sin365 + cos365, random = ~ 1 | ID,
  family = poisson, data = d,
  correlation=nlme::corAR1(form=~dayOfSeries|ID))
```


```
## iteration 1
```

```
summary(fit)
```

```
## Linear mixed-effects model fit by maximum likelihood
## Data: d
## AIC BIC logLik
## NA NA NA
##
## Random effects:
## Formula: ~1 | ID
## (Intercept) Residual
## StdDev: 1.149069e-05 0.841689
##
## Correlation Structure: AR(1)
## Formula: ~dayOfSeries | ID
## Parameter estimate(s):
## Phi
## 0.4926123
## Variance function:
```


```
## Structure: fixed weights
## Formula: ~invwt
## Fixed effects: y ~ yearMinus2000 + sin365 + cos365
##           Value   Std.Error   DF   t-value p-value
## (Intercept)  0.7980540 0.015203158 6936  52.49265    0
## yearMinus2000 0.0991582 0.001070583 6936  92.62077    0
## sin365        1.4074339 0.010596649 6936 132.81876    0
## cos365       -0.5389807 0.008876447 6936 -60.72031    0
## Correlation:
##           (Intr) yM2000 sin365
## yearMinus2000 -0.832
## sin365        -0.409  0.000
## cos365         0.186  0.000 -0.158
##
## Standardized Within-Group Residuals:
##           Min           Q1           Med           Q3           Max
## -2.89886753 -0.75775062 -0.05982255  0.60730690  6.49964494
##
## Number of Observations: 6940
## Number of Groups: 1
```

```
pacf(residuals(fit, type = "normalized")) # this is for AR
```



```
acf(residuals(fit, type = "normalized")) # this is for MA
```


Series residuals(fit, type = "normalized")



```

b1 <- 1.3936185 # sin coefficient
b2 <- -0.5233866 # cos coefficient
amplitude <- sqrt(b1^2 + b2^2)
p <- atan(b1/b2) * 365/2/pi
if (p > 0) {
  peak <- p
  trough <- p + 365/2
} else {
  peak <- p + 365/2
  trough <- p + 365
}
if (b1 < 0) {
  g <- peak
  peak <- trough
  trough <- g
}
print(sprintf("amplitude is estimated as %s, peak is estimated as %s, trough is estimated as %s",round(

```

```
## [1] "amplitude is estimated as 1.49, peak is estimated as 112, trough is estimated as 295"
```


```
print(sprintf("true values are: amplitude: %s, peak: %s, trough: %s",round(AMPLITUDE,2),round(365/4+SEA

```


```
## [1] "true values are: amplitude: 1.5, peak: 111, trough: 294"
```

4.1 Showing that `tscount::tsglm` gets the same results as `MASS::glmPQL`

```
library(MASS)
correlatedError <- as.numeric(arima.sim(model=list("ar"=c(-0.9)), n=1000, rand.gen = rnorm))
pacf(correlatedError) # this is for AR
```



```
acf(correlatedError) # this is for MA
```



```
d <- data.frame(correlatedError)
d$independentError <- rnorm(nrow(d))
d$x <- rnorm(nrow(d))
d$yCorrelated <- 2*d$x+d$correlatedError
d$yIndependent <- 2*d$x+d$independentError
d$ID <- 1
d$time <- 1:nrow(d)
```


```
summary(lm(yIndependent~x,data=d))
```

```
##
## Call:
## lm(formula = yIndependent ~ x, data = d)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3.7055 -0.6919  0.0129  0.6823  2.8945
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.01133    0.03139  -0.361   0.718
## x             2.06473    0.03103  66.541 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.9922 on 998 degrees of freedom
```


```
## Multiple R-squared:  0.8161, Adjusted R-squared:  0.8159
## F-statistic:  4428 on 1 and 998 DF,  p-value: < 2.2e-16
summary(fit <- lm(yCorrelated~x,data=d))

##
## Call:
## lm(formula = yCorrelated ~ x, data = d)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -7.8825 -1.3737 -0.0576  1.4968  7.4766
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.03095    0.07241  -0.427   0.669
## x            1.95682    0.07158  27.337 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.289 on 998 degrees of freedom
## Multiple R-squared:  0.4282, Adjusted R-squared:  0.4276
## F-statistic: 747.3 on 1 and 998 DF,  p-value: < 2.2e-16
pacf(residuals(fit)) # this is for AR
```

Series residuals(fit)



```
acf(residuals(fit)) # this is for MA
```



```
# independent data, no correlation structure needed
fit <- MASS::glmmpQL(yIndependent ~ x, random = ~ 1 | ID,
  family = gaussian, data = d,
  correlation=nlme::corAR1())
```

```
## iteration 1
```

```
## iteration 2
```

```
summary(fit)
```

```
## Linear mixed-effects model fit by maximum likelihood
```

```
## Data: d
```

```
## AIC BIC logLik
```

```
## NA NA NA
```

```
##
```

```
## Random effects:
```

```
## Formula: ~1 | ID
```

```
## (Intercept) Residual
```

```
## StdDev: 3.140954e-05 0.9912063
```

```
##
```

```
## Correlation Structure: AR(1)
```

```
## Formula: ~1 | ID
```

```
## Parameter estimate(s):
```


```
## Phi
```

```
## 0.04573347
```


```
## Variance function:
```

```
## Structure: fixed weights
```


```
## Formula: ~invwt
## Fixed effects: yIndependent ~ x
##           Value Std.Error DF   t-value p-value
## (Intercept) -0.0112317 0.03285689 998 -0.34184  0.7325
## x           2.0682689 0.03099327 998 66.73283  0.0000
## Correlation:
##   (Intr)
## x 0.028
##
## Standardized Within-Group Residuals:
##           Min           Q1           Med           Q3           Max
## -3.74008252 -0.69738448  0.01047609  0.68729125  2.92537677
##
## Number of Observations: 1000
## Number of Groups: 1
pacf(residuals(fit, type = "response")) # this is for AR
```




```
acf(residuals(fit, type = "response")) # this is for MA
```



```
pacf(residuals(fit, type = "normalized")) # this is for AR
```



```
acf(residuals(fit, type = "normalized")) # this is for MA
```


Series residuals(fit, type = "normalized")


```
# dependent data, needs correlation structure, no correlation structure
fit <- MASS::glmmpQL(yCorrelated ~ x, random = ~ 1 | ID,
  family = gaussian, data = d)
```

```
## iteration 1
```


```
summary(fit)
```

```
## Linear mixed-effects model fit by maximum likelihood
## Data: d
##   AIC BIC logLik
##   NA  NA    NA
##
## Random effects:
## Formula: ~1 | ID
##      (Intercept) Residual
## StdDev: 7.486629e-05 2.286563
##
## Variance function:
## Structure: fixed weights
## Formula: ~invwt
## Fixed effects: yCorrelated ~ x
##              Value Std.Error DF   t-value p-value
## (Intercept) -0.0309491 0.07241167 998 -0.427404  0.6692
## x             1.9568162 0.07158024 998 27.337380  0.0000
## Correlation:
## (Intr)
```


```
## x 0.03
##
## Standardized Within-Group Residuals:
##      Min      Q1      Med      Q3      Max
## -3.44731225 -0.60078114 -0.02520784  0.65459886  3.26978586
##
## Number of Observations: 1000
## Number of Groups: 1
pacf(residuals(fit, type = "response")) # this is for AR
```




```
acf(residuals(fit, type = "response")) # this is for MA
```



```
pacf(residuals(fit, type = "normalized")) # this is for AR
```



```
acf(residuals(fit, type = "normalized")) # this is for MA
```




```
# dependent data, correct correlation structure
fit <- MASS::glmmpQL(yCorrelated ~ x, random = ~ 1 | ID,
  family = gaussian, data = d,
  correlation=nlme::corAR1())
```

```
## iteration 1
## iteration 2
```


```
summary(fit)
```

```
## Linear mixed-effects model fit by maximum likelihood
## Data: d
##   AIC BIC logLik
##   NA  NA    NA
##
## Random effects:
## Formula: ~1 | ID
##      (Intercept) Residual
## StdDev: 1.643326e-05 2.282467
##
## Correlation Structure: AR(1)
## Formula: ~1 | ID
## Parameter estimate(s):
##      Phi
## -0.9031247
## Variance function:
## Structure: fixed weights
```


```
## Formula: ~invwt
## Fixed effects: yCorrelated ~ x
##           Value Std.Error DF t-value p-value
## (Intercept) -0.0296893 0.01632325 998 -1.81883 0.0692
## x           2.0072905 0.02255834 998 88.98217 0.0000
## Correlation:
## (Intr)
## x 0.042
##
## Standardized Within-Group Residuals:
##           Min           Q1           Med           Q3           Max
## -3.46470223 -0.60212874 -0.01867443  0.66118043  3.29397505
##
## Number of Observations: 1000
## Number of Groups: 1
pacf(residuals(fit, type = "response")) # this is for AR
```




```
acf(residuals(fit, type = "response")) # this is for MA
```



```
pacf(residuals(fit, type = "normalized")) # this is for AR
```




```
acf(residuals(fit, type = "normalized")) # this is for MA
```


```
library(MASS)

bacteria$x <- 1
fit <- glmmpQL(as.numeric(y) ~ trt + I(week > 2), random = ~ 1 | x,
               family = poisson, data = bacteria)
```

```
## iteration 1
acf(residuals(fit,type="normalized"))
```




```
pacf(residuals(fit,type="normalized"))
```




```
bacteria$x <- 1
fit <- glmmPQL(as.numeric(y) ~ trt + I(week > 2), random = ~ 1 | x,
               family = poisson, data = bacteria,
               correlation=nlme::corAR1())
```

```
## iteration 1
```

```
## iteration 2
acf(residuals(fit,type="normalized"))
```




```
pacf(residuals(fit,type="normalized"))
```




```
fit <- glmmPQL(as.numeric(y) ~ trt + I(week > 2), random = ~ 1 | ID,
               family = poisson, data = bacteria)
```

```
## iteration 1
## iteration 2
```

```
## iteration 3
acf(residuals(fit,type="normalized"))
```



```
pacf(residuals(fit,type="normalized"))
```



```
fit <- glmmPQL(as.numeric(y) ~ trt + I(week > 2), random = ~ 1 | ID,
               family = poisson, data = bacteria,
               correlation=nlme::corAR1(form=~ 1 | ID))
```


```
## iteration 1
```

```
## iteration 2
```


```
## iteration 3
```

```
## iteration 4
```

```
acf(residuals(fit,type="normalized"))
```



```
pacf(residuals(fit,type="normalized"))
```

Chapter 5

Variables

```
library(data.table)
library(lme4)

## Loading required package: Matrix
## Loading required package: methods
set.seed(4)

fylkeIntercepts <- data.table(fylke=1:20,fylkeIntercepts=rnorm(20))

d <- data.table(fylke=rep(1:20,each=100))
d <- merge(d,fylkeIntercepts,by="fylke")
d[,mainIntercept:=3]
d[,x:=runif(.N)]
d[,mu := exp(mainIntercept + fylkeIntercepts + 3*x)]
d[,y:=rpois(.N,mu)]

summary(fit <- lme4::glmer(y~x + (1|fylke),data=d,family=poisson()))

## Generalized linear mixed model fit by maximum likelihood (Laplace
##   Approximation) [glmerMod]
##   Family: poisson ( log )
## Formula: y ~ x + (1 | fylke)
##   Data: d
##
##           AIC          BIC    logLik deviance df.resid
## 15508.5 15525.3 -7751.3 15502.5      1997
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -3.1132 -0.6422 -0.0260  0.6556  3.6029
##
## Random effects:
##   Groups Name            Variance Std.Dev.
##   fylke  (Intercept) 0.6167   0.7853
## Number of obs: 2000, groups: fylke, 20
##
## Fixed effects:
```

```
##           Estimate Std. Error z value Pr(>|z|)
## (Intercept) 3.378858  0.175660   19.2  <2e-16 ***
## x           2.990811  0.005991  499.2  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##   (Intr)
## x -0.024
```

Chapter 6

Variables

```
library(data.table)
library(ggplot2)
set.seed(4)

AMPLITUDE <- 1.5
SEASONAL_HORIZONTAL_SHIFT <- 20

fylkeIntercepts <- data.table(fylke=1:20,fylkeIntercepts=rnorm(20))

d <- data.table(date=seq.Date(
  from=as.Date("2010-01-01"),
  to=as.Date("2015-12-31"),
  by=1))
d[,year:=as.numeric(format.Date(date,"%G"))]
d[,week:=as.numeric(format.Date(date,"%V"))]
d[,month:=as.numeric(format.Date(date,"%m"))]

temp <- vector("list",length=20)
for(i in 1:20){
  temp[[i]] <- copy(d)
  temp[[i]][,fylke:=i]
}
d <- rbindlist(temp)

d[,yearMinus2000:=year-2000]
d[,dayOfSeries:=1:.N]

d[,dayOfYear:=as.numeric(format.Date(date,"%j"))]
d[,seasonalEffect:=sin(2*pi*(dayOfYear-SEASONAL_HORIZONTAL_SHIFT)/365)]
d[,mu := exp(0.1 + yearMinus2000*0.1 + seasonalEffect*AMPLITUDE)]
d[,y:=rpois(.N,mu)]
#d[,y:=round(as.numeric(arima.sim(model=list("ar"=c(0.5)), rand.gen = rpois, n=nrow(d), lambda=mu)))]
```

We then drill down into a few years, and see a clear seasonal trend

```
q <- ggplot(d[fylke==1],aes(x=dayOfYear,y=y))
q <- q + facet_wrap(~year)
q <- q + geom_point()
```

```

q <- q + stat_smooth(colour="red")
q

## `geom_smooth()` using method = 'loess'

## Warning in simpleLoess(y, x, w, span, degree = degree, parametric =
## parametric, : span too small. fewer data values than degrees of freedom.

## Warning in simpleLoess(y, x, w, span, degree = degree, parametric =
## parametric, : pseudoinverse used at 0.99

## Warning in simpleLoess(y, x, w, span, degree = degree, parametric =
## parametric, : neighborhood radius 1.01

## Warning in simpleLoess(y, x, w, span, degree = degree, parametric =
## parametric, : reciprocal condition number 0

## Warning in simpleLoess(y, x, w, span, degree = degree, parametric =
## parametric, : There are other near singularities as well. 1.0201

## Warning in predLoess(object$y, object$x, newx = if
## (is.null(newdata)) object$x else if (is.data.frame(newdata))
## as.matrix(model.frame(delete.response(terms(object))), : span too small.
## fewer data values than degrees of freedom.

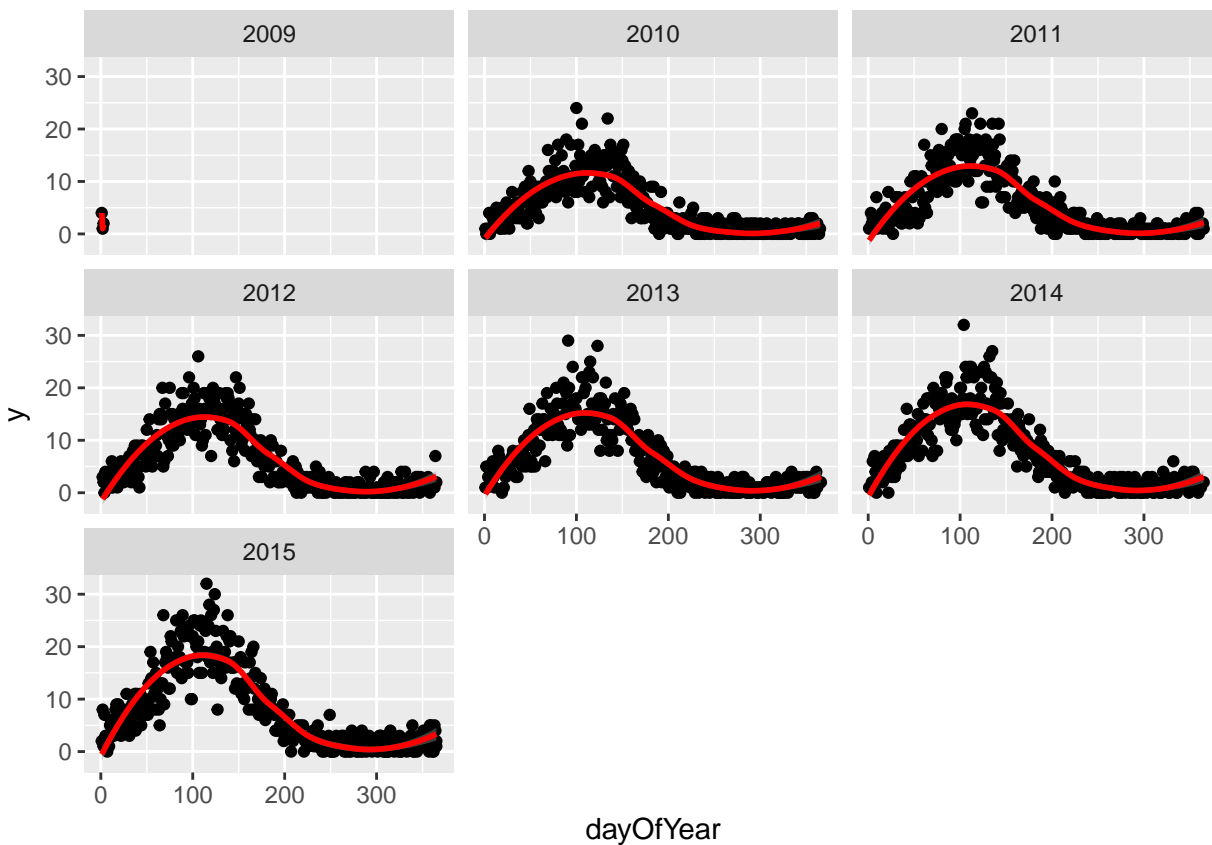
## Warning in predLoess(object$y, object$x, newx = if
## (is.null(newdata)) object$x else if (is.data.frame(newdata))
## as.matrix(model.frame(delete.response(terms(object))), : pseudoinverse used
## at 0.99

## Warning in predLoess(object$y, object$x, newx = if
## (is.null(newdata)) object$x else if (is.data.frame(newdata))
## as.matrix(model.frame(delete.response(terms(object))), : neighborhood radius
## 1.01

## Warning in predLoess(object$y, object$x, newx = if
## (is.null(newdata)) object$x else if (is.data.frame(newdata))
## as.matrix(model.frame(delete.response(terms(object))), : reciprocal
## condition number 0


## Warning in predLoess(object$y, object$x, newx = if
## (is.null(newdata)) object$x else if (is.data.frame(newdata))
## as.matrix(model.frame(delete.response(terms(object))), : There are other
## near singularities as well. 1.0201

```



The Lomb-Scargle Periodogram shows a clear seasonality with a period of 365 days

```
lomb::lsp(d$y, from=100, to=500, ofac=1, type="period")
```



```
d[,cos365:=cos(dayOfYear*2*pi/365)]
d[,sin365:=sin(dayOfYear*2*pi/365)]

fit <- MASS::glmmPQL(y~yearMinus2000+sin365 + cos365, random = ~ 1 | fylke,
  family = poisson, data = d,
  correlation=nlme::corAR1(form=~dayOfSeries|fylke))


## iteration 1
summary(fit)

## Linear mixed-effects model fit by maximum likelihood
## Data: d
## AIC BIC logLik
## NA NA NA
##
## Random effects:
## Formula: ~1 | fylke
## (Intercept) Residual
## StdDev: 1.708256e-05 0.9976713
##
## Correlation Structure: AR(1)
## Formula: ~dayOfSeries | fylke
## Parameter estimate(s):
## Phi
## 0.002841665
## Variance function:
```




```
## Structure: fixed weights
## Formula: ~invwt
## Fixed effects: y ~ yearMinus2000 + sin365 + cos365
##           Value   Std.Error    DF   t-value p-value
## (Intercept)  0.1122528 0.014529606 43797    7.7258    0
## yearMinus2000 0.0989047 0.001112632 43797   88.8926    0
## sin365        1.4095094 0.003705852 43797  380.3469    0
## cos365       -0.5109372 0.003092449 43797 -165.2209    0
## Correlation:
##           (Intr) yM2000 sin365
## yearMinus2000 -0.979
## sin365        -0.150  0.000
## cos365         0.065 -0.001 -0.151
##
## Standardized Within-Group Residuals:
##           Min      Q1      Med      Q3      Max
## -3.1968230 -0.8238741 -0.0750183  0.6340046  5.8245241
##
## Number of Observations: 43820
## Number of Groups: 20
```

```
pacf(residuals(fit, type = "normalized")) # this is for AR
```



```
acf(residuals(fit, type = "normalized")) # this is for MA
```



```

b1 <- 1.4007640 # sin coefficient
b2 <- -0.5234863 # cos coefficient
amplitude <- sqrt(b1^2 + b2^2)
p <- atan(b1/b2) * 365/2/pi
if (p > 0) {
  peak <- p
  trough <- p + 365/2
} else {
  peak <- p + 365/2
  trough <- p + 365
}
if (b1 < 0) {
  g <- peak
  peak <- trough
  trough <- g
}
print(sprintf("amplitude is estimated as %s, peak is estimated as %s, trough is estimated as %s",round(
## [1] "amplitude is estimated as 1.5, peak is estimated as 112, trough is estimated as 295"
print(sprintf("true values are: amplitude: %s, peak: %s, trough: %s",round(AMPLITUDE,2),round(365/4+SEA
## [1] "true values are: amplitude: 1.5, peak: 111, trough: 294"

```

Chapter 7

Variables

```
library(data.table)
library(ggplot2)
set.seed(4)

AMPLITUDE <- 1.5
SEASONAL_HORIZONTAL_SHIFT <- 20

fylkeIntercepts <- data.table(fylke=1:20,fylkeIntercepts=rnorm(20))

d <- data.table(date=seq.Date(
  from=as.Date("2010-01-01"),
  to=as.Date("2015-12-31"),
  by=1))
d[,year:=as.numeric(format.Date(date,"%G"))]
d[,week:=as.numeric(format.Date(date,"%V"))]
d[,month:=as.numeric(format.Date(date,"%m"))]

temp <- vector("list",length=20)
for(i in 1:20){
  temp[[i]] <- copy(d)
  temp[[i]][,fylke:=i]
}
d <- rbindlist(temp)

d[,yearMinus2000:=year-2000]
d[,dayOfSeries:=1:.N]

d[,dayOfYear:=as.numeric(format.Date(date,"%j"))]
d[,seasonalEffect:=sin(2*pi*(dayOfYear-SEASONAL_HORIZONTAL_SHIFT)/365)]
d[,mu := exp(0.1 + yearMinus2000*0.1 + seasonalEffect*AMPLITUDE)]
d[,y:=rpois(.N,mu)]
d[,y:=round(as.numeric(arima.sim(model=list("ar"=c(0.5)), rand.gen = rpois, n=nrow(d), lambda=mu)))]
```

We then drill down into a few years, and see a clear seasonal trend

```
q <- ggplot(d[fylke==1],aes(x=dayOfYear,y=y))
q <- q + facet_wrap(~year)
q <- q + geom_point()
```

```

q <- q + stat_smooth(colour="red")
q

## `geom_smooth()` using method = 'loess'

## Warning in simpleLoess(y, x, w, span, degree = degree, parametric =
## parametric, : span too small. fewer data values than degrees of freedom.

## Warning in simpleLoess(y, x, w, span, degree = degree, parametric =
## parametric, : pseudoinverse used at 0.99

## Warning in simpleLoess(y, x, w, span, degree = degree, parametric =
## parametric, : neighborhood radius 1.01

## Warning in simpleLoess(y, x, w, span, degree = degree, parametric =
## parametric, : reciprocal condition number 0

## Warning in simpleLoess(y, x, w, span, degree = degree, parametric =
## parametric, : There are other near singularities as well. 1.0201

## Warning in predLoess(object$y, object$x, newx = if
## (is.null(newdata)) object$x else if (is.data.frame(newdata))
## as.matrix(model.frame(delete.response(terms(object))), : span too small.
## fewer data values than degrees of freedom.

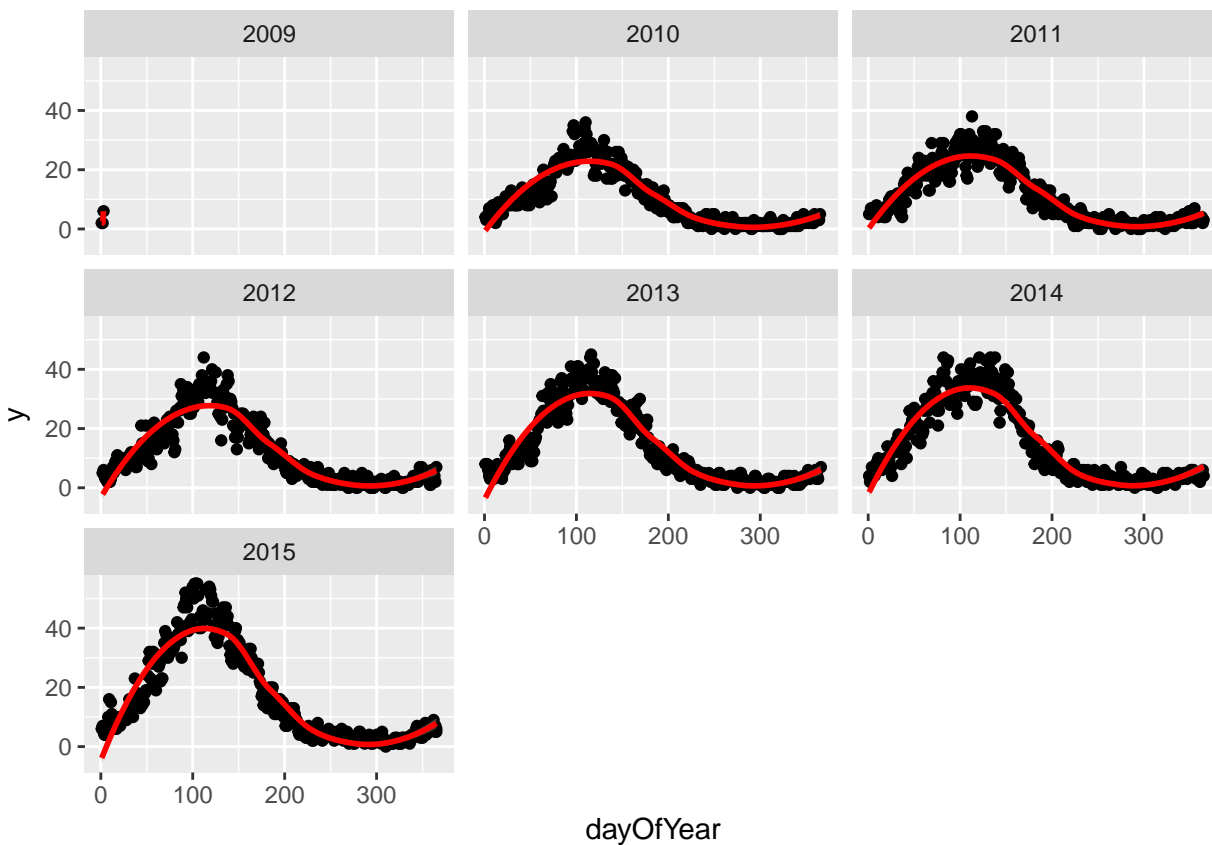
## Warning in predLoess(object$y, object$x, newx = if
## (is.null(newdata)) object$x else if (is.data.frame(newdata))
## as.matrix(model.frame(delete.response(terms(object))), : pseudoinverse used
## at 0.99

## Warning in predLoess(object$y, object$x, newx = if
## (is.null(newdata)) object$x else if (is.data.frame(newdata))
## as.matrix(model.frame(delete.response(terms(object))), : neighborhood radius
## 1.01

## Warning in predLoess(object$y, object$x, newx = if
## (is.null(newdata)) object$x else if (is.data.frame(newdata))
## as.matrix(model.frame(delete.response(terms(object))), : reciprocal
## condition number 0

## Warning in predLoess(object$y, object$x, newx = if
## (is.null(newdata)) object$x else if (is.data.frame(newdata))
## as.matrix(model.frame(delete.response(terms(object))), : There are other
## near singularities as well. 1.0201


```



The Lomb-Scargle Periodogram shows a clear seasonality with a period of 365 days

```
lomb::lsp(d$y, from=100, to=500, ofac=1, type="period")
```

Lomb–Scargle Periodogram



```
d[,cos365:=cos(dayOfYear*2*pi/365)]
d[,sin365:=sin(dayOfYear*2*pi/365)]


fit <- MASS::glmmPQL(y~yearMinus2000+sin365 + cos365, random = ~ 1 | fylke,
  family = poisson, data = d,
  correlation=nlme::corAR1(form=~dayOfSeries|fylke))

## iteration 1
summary(fit)


## Linear mixed-effects model fit by maximum likelihood
## Data: d
## AIC BIC logLik
## NA NA NA
##
## Random effects:
## Formula: ~1 | fylke
## (Intercept) Residual
## StdDev: 0.003916838 0.822866
##
## Correlation Structure: AR(1)
## Formula: ~dayOfSeries | fylke
## Parameter estimate(s):
## Phi
## 0.4771948
## Variance function:
```

```
## Structure: fixed weights
## Formula: ~invwt
## Fixed effects: y ~ yearMinus2000 + sin365 + cos365
##           Value   Std.Error   DF   t-value p-value
## (Intercept)  0.8181967 0.014201621 43797   57.6129    0
## yearMinus2000 0.0982444 0.001085637 43797   90.4947    0
## sin365        1.4007640 0.003607254 43797  388.3187    0
## cos365       -0.5234863 0.003020395 43797 -173.3171    0
## Correlation:
##           (Intr) yM2000 sin365
## yearMinus2000 -0.977
## sin365        -0.149  0.001
## cos365         0.067 -0.001 -0.153
##
## Standardized Within-Group Residuals:
##           Min           Q1           Med           Q3           Max
## -3.40197340 -0.70762882 -0.06465241  0.62676531  5.47900204
##
## Number of Observations: 43820
## Number of Groups: 20
```

```
pacf(residuals(fit, type = "normalized")) # this is for AR
```



```
acf(residuals(fit, type = "normalized")) # this is for MA
```



```

b1 <- 1.4007640 # sin coefficient
b2 <- -0.5234863 # cos coefficient
amplitude <- sqrt(b1^2 + b2^2)
p <- atan(b1/b2) * 365/2/pi
if (p > 0) {
  peak <- p
  trough <- p + 365/2
} else {
  peak <- p + 365/2
  trough <- p + 365
}
if (b1 < 0) {
  g <- peak
  peak <- trough
  trough <- g
}
print(sprintf("amplitude is estimated as %s, peak is estimated as %s, trough is estimated as %s",round(
## [1] "amplitude is estimated as 1.5, peak is estimated as 112, trough is estimated as 295"
print(sprintf("true values are: amplitude: %s, peak: %s, trough: %s",round(AMPLITUDE,2),round(365/4+SEA
## [1] "true values are: amplitude: 1.5, peak: 111, trough: 294"

```